


Is it time to replace mmap?


A history of virtual address management
(and a proposal for the future)

Brooks Davis
SRI International
BSDTW 2017. Taipei, Taiwan

Memory



A bit of computer history



ENIAC c.1945



EDSAC c.1949



PDP-11 c.1970

I386 1985




Photo sources:

ENIAC: Two women operating ENIAC - U.S. Army Photo

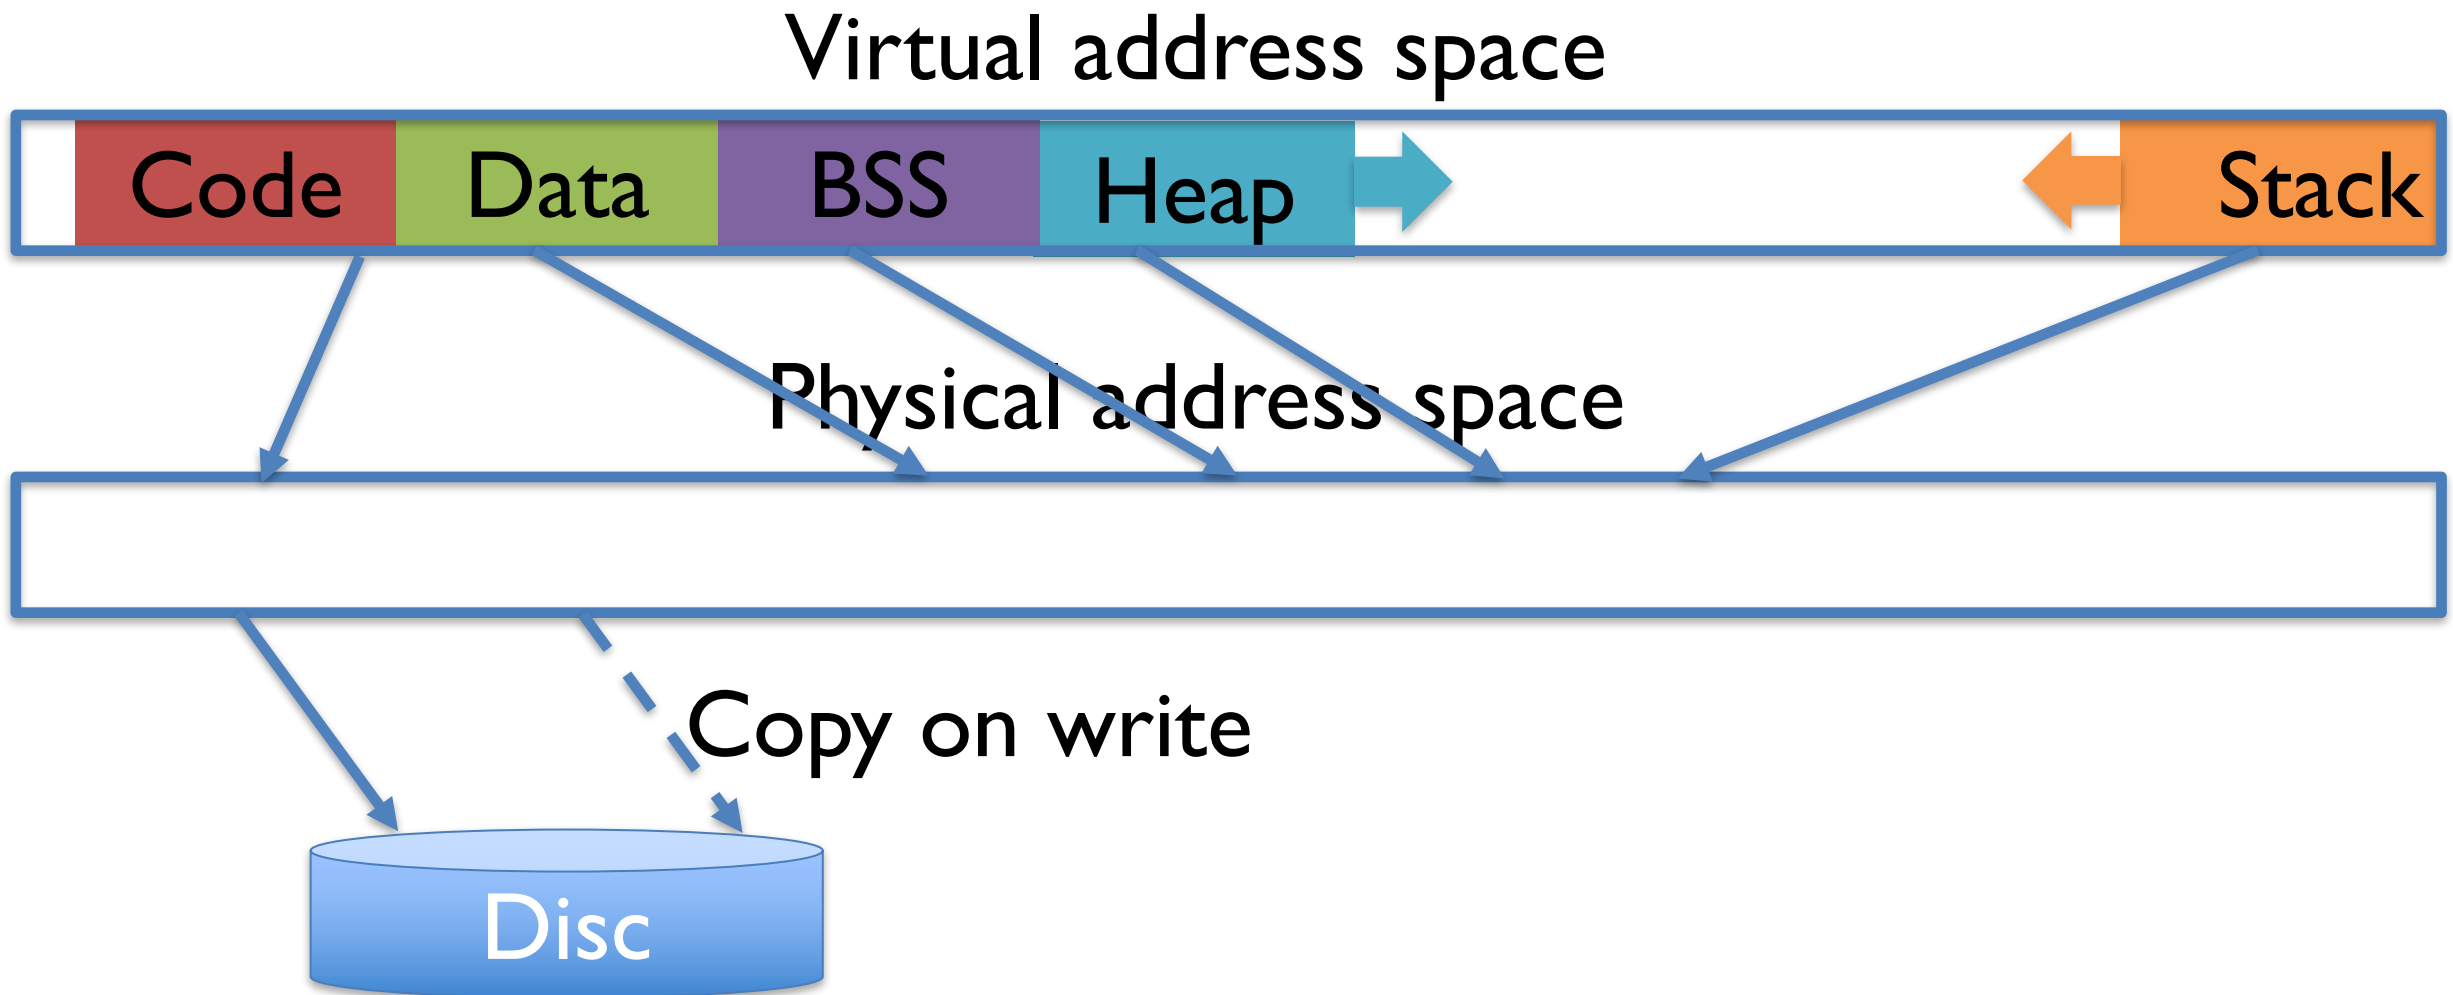
EDSAC: EDSAC I, R.Hill operating - Copyright Computer Laboratory, University of Cambridge.

Reproduced by permission. PDP-11: DEC - PDP-11 - Ken Thompson and Dennis Ritchie - Courtesy Computer History Museum


Process address space




Process address space



Process address space



UNIX and BSD




break.2 (V3 Unix)

break sets the system's idea of the highest location used by the program to addr.

Locations greater than addr and below the stack pointer are **not swapped and are thus liable to unexpected modification.**

UNIX and BSD



break.2 (V4 Unix)

Break sets the system's idea of the lowest location not used by the program to *addr* (rounded up to the next multiple of 64 bytes).

Locations not less than *addr* and below the stack pointer are not in the address space and will thus cause a memory violation if accessed.

break.2 (V4 Unix) (cont)

```
char *sbrk(incr)
```

...

From C, the calling sequence is different; *incr* more bytes are added to the program's data space and a pointer to the start of the new area is returned.

break.2 (V4 Unix) (cont)

When a program begins execution via `exec` the break is set at the highest location defined by the program and data storage areas. Ordinarily, therefore, only programs with growing data areas need to use *break*.

UNIX and BSD

1973 V4:

sbreak

system call

now provides
protection.

sbrk() introduced

1975 V5:

brk() introduced

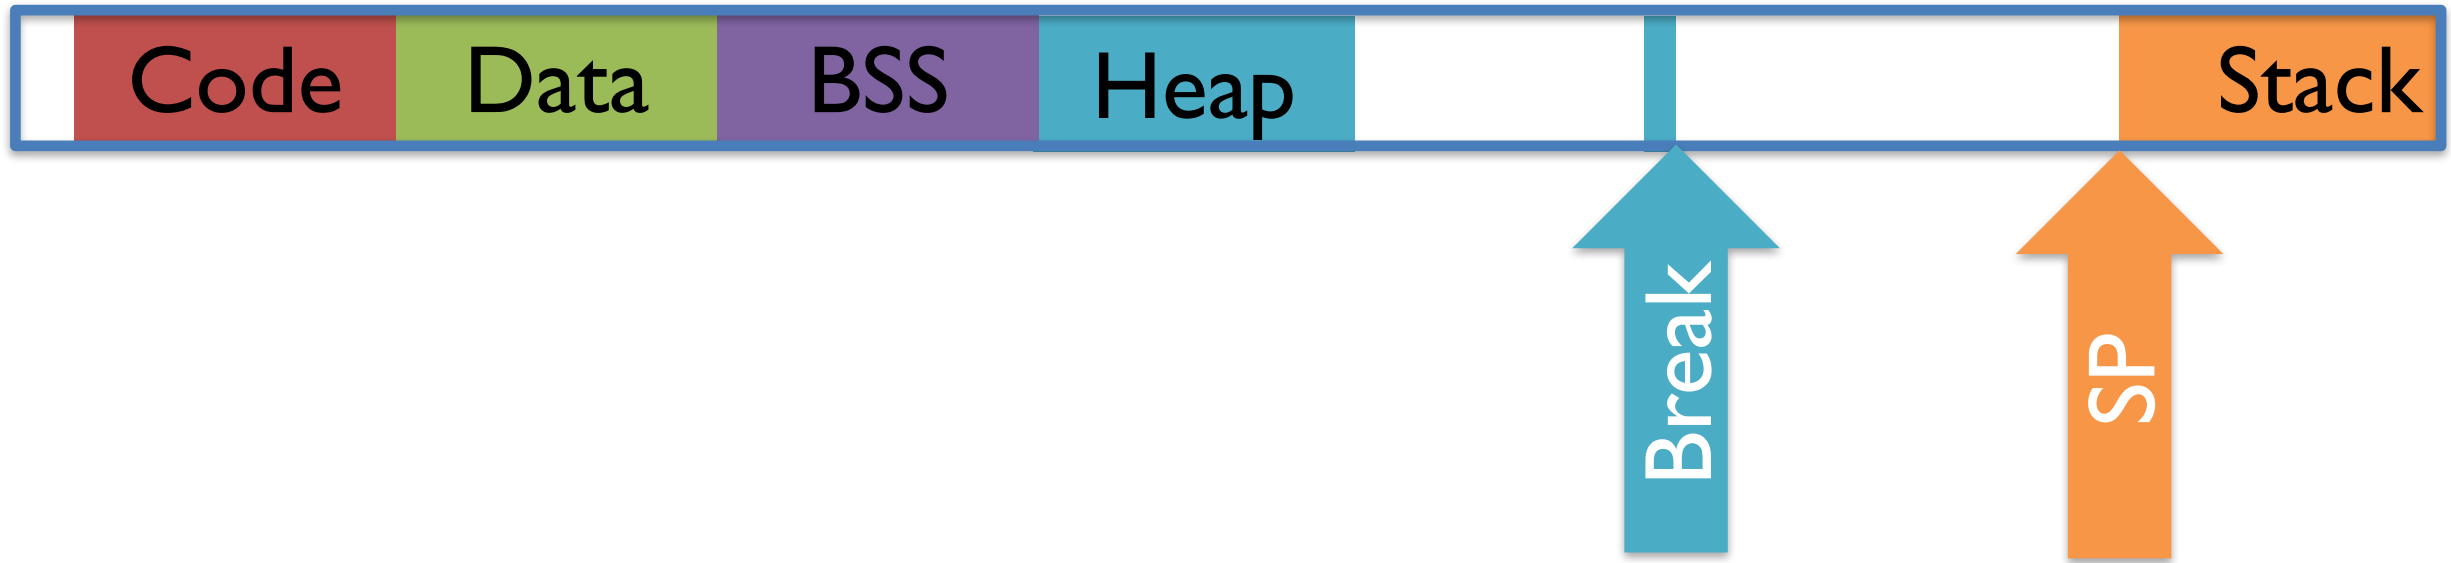
1983

4.2BSD:

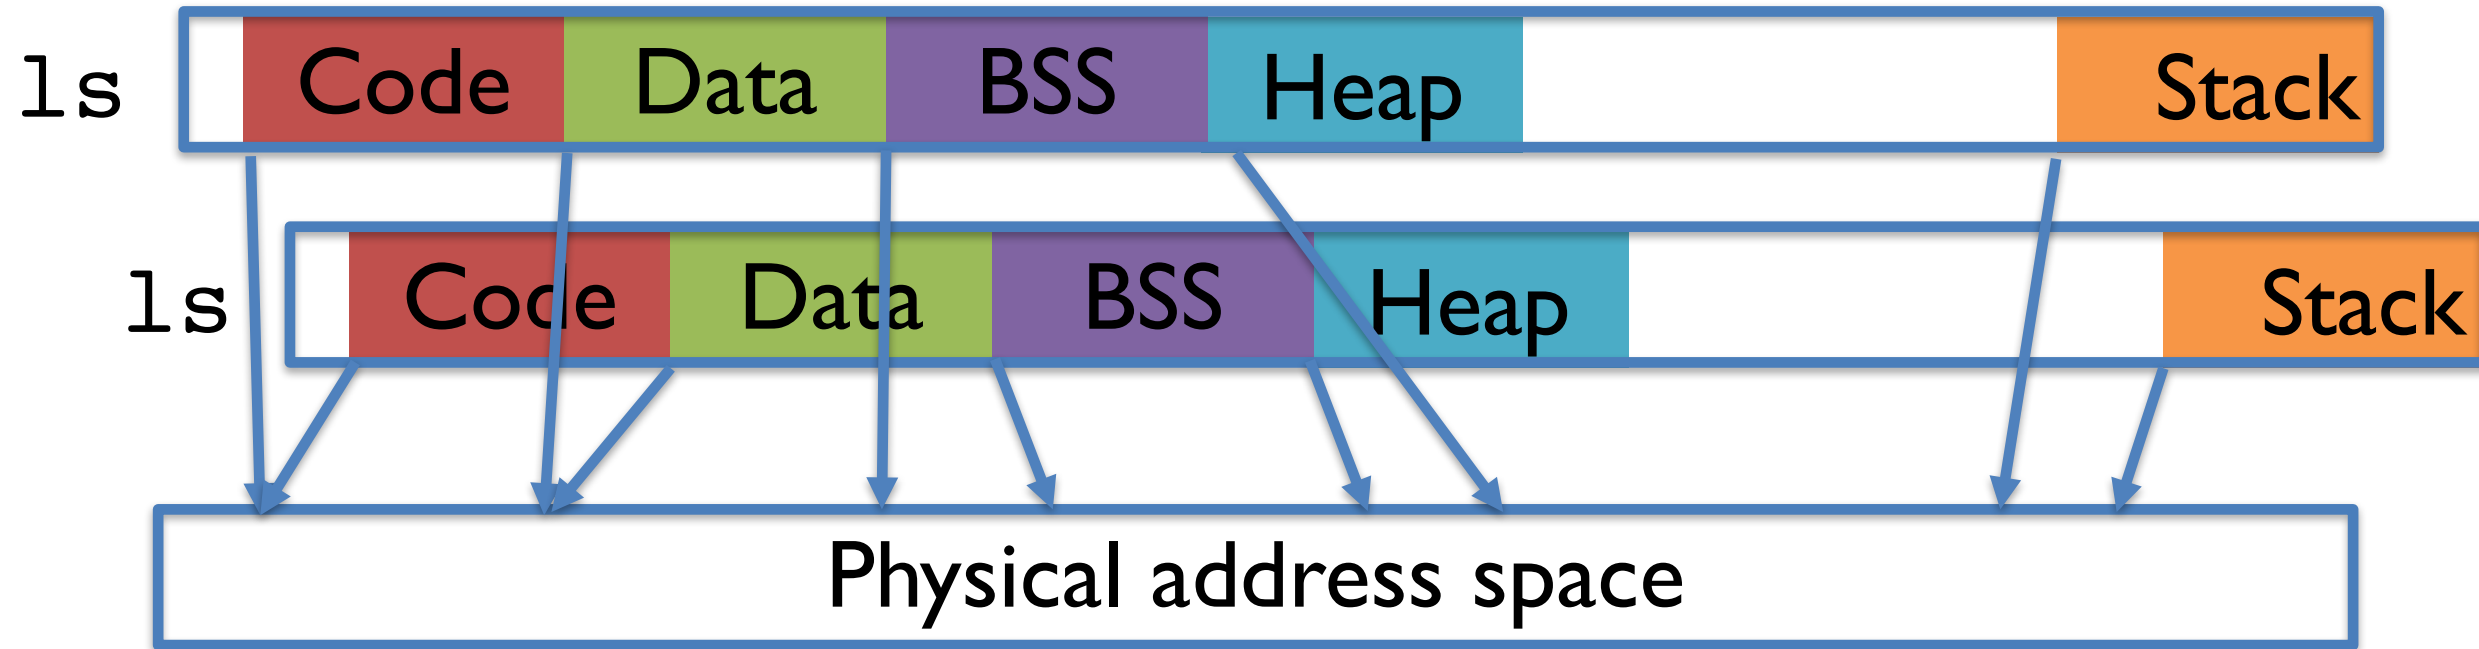
First
references
to mmap()




Heap fragmentation



Memory sharing



Dynamic linking



Multi-threaded programs

With sbrk() only



Ideal



4.2BSD memory interfaces

- `mmap()`
 - Allocate address space
 - Alter backing mappings
- `mremap()`
 - Relocate or extend mapping
- `munmap()`
 - Remove backing

4.2BSD memory interfaces

- `mprotect()`
 - Alter page protections
- `madvise()`
 - Hint usage to kernel
- `mincore()`
 - Query backing status
- `sbrk()`
 - Extend or reduce “break”
- `sstk()`
 - Extend or reduce stack

4.2BSD memory interfaces

- `mprotect()`
 - Alter page protection
- `madvise()`
- `sbrk()`
 - Extend or reduce stack
- `brk()`
- `setbrk()`
- `query backing status`

Only `sbrk()` implemented!

UNIX and BSD

1973 V4:

sbreak

system call

now provides
protection.

sbrk() introduced

1975 V5:

brk() introduced

1983

4.2BSD:

First
references
to mmap()

1990 4.3-Reno:

mmap()

implemented
with VM from
Mach




UNIX and BSD

1990 4.3-Reno:

mmap()
implemented
with VM from
Mach

2003 OpenBSD 3.3:
Implements W^X



W^X and JITs

- Prohibits pages from having both PROT_EXEC and PROT_WRITE simultaneously
- JITs need to write then execute!
- Solution: Map PROT_WRITE then remove PROT_WRITE and add PROT_EXEC
- New problem: most pages should not become executable, but mmap() cannot express this!

UNIX and BSD

1990 4.3-Reno:

mmap()
implemented
with VM from
Mach

2012 CHERI Project

2003 OpenBSD 3.3:
Implements W^X



CHERI pointers

- Pointers with bounds and permissions
 - With strong monotonicity guarantees
- Want W^X for pointers (in addition to pages)
- API changes required:
 - Should `mprotect()` return a pointer?
 - Should some other mechanism be used?

mmap() functionality issues

- Interface conflates address reservation and mapping
 - Lack of boundaries between reservations leads to bugs: e.g. Stack Clash
- Lack of expressiveness
 - No portable way to express alignment
 - No way to express maximum permission

mmap() API issues

- Too many arguments
 - Can you remember them all?
 - Many calls don't use them all
- Too many failure modes:
 - FreeBSD 11: 19 documented errors (15 use the same error code)

Other mmap() issues

- No support for mapping more pages than requested
 - Can't round up to superpage size
 - CHERI bounds compression requires rounding for very large allocations
- No concept of address space ownership
 - Math errors mean changing the wrong region

RFC: mmap (1/3)

- `int mreserve(cm_t *handlep, size_t length, vaddr_t hint, int prot, cmreq_t *cmr);`
 - Reserve a region, optionally mapping.
- `int mgetptr(cm_t handle, void **ptrp);`
 - Get a pointer to the region.

RFC: mmap (2/3)

- `int mmap(cm_t handle, cmreq_t *cmr);`
 - Replace (part of) a region's mappings.
- `int mclose(cm_t handle);`
 - Close a handle.
- `int mrestrict(cm_t handle, XX ops, XX *oops);`
 - Restrict the set of operations on a handle

RFC: mmap (2/3)

- `int cmstat(cm_t handle, size_t index, struct cm_stat * cs)`
 - Return data on a series of submaps
- `cmadvise()`, `cmincore()`, `cminherit()`, `cmsync()`, `cmunmap()`
 - Like `mmap()` counterparts, but within region

More on map requests

- Request object rather than many arguments
 - `cm_request_t` following `pthread_attr_t` model
- Accessor functions to set up request
- Goal: useful defaults
 - Ideally, requests should always be valid

CHERI extensions

- `int cmgetcap(cm_t cookie, void **ptrp, perm_t perms)`
 - Get a capability pointer
- `int cmandperm(cm_t cookie, perm_t perms, perm_t *operms)`
 - Reduce the set of allowed permissions

Should we replace mmap()?

Yes or No?

ON MORETHING

UNIX and BSD


1990 4.3-Reno:

mmap()
implemented
with VM from
Mach

2012 CHERI Project

2003 OpenBSD 3.3:
Implements W^X

2016 FreeBSD 11:
No sbrk() on Arm64
and RISC-V



Removing sbrk()

- Mostly incorrect attempts to measure heap use
 - Usually can be disabled, but some force required
- A few internal allocators
 - Usually can be disabled
- Some LISP interpreters
 - Mostly unpopular ones

Removing sbrk() (cont.)



Adrian Chadd

@erikarn

Follow



TIL: Best way to erm, "win" the editor wars is to ship a new platform with ^{out} sbrk support.
[#freebsd](#) did that on arm64 - and no emacs!

5:17 PM - 18 Nov 2016